

# Machine Learning Lecture 6

## Logistic Regression (Linear Classifiers)

Dr. Ioannis Patras

Slides thanks: Tim Hospedales

## Course Context

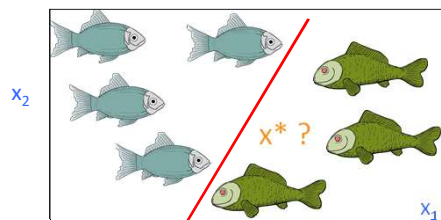
- Supervised Learning
  - (Linear) regression
  - (Linear) Classifiers and Logistic Regression
  - Neural Networks
- Unsupervised
  - Clustering
  - Density Estimation
  - HMMs

## Supervised Learning

- Applications where the training data comprises examples of input vectors along with corresponding target vectors
- **Regression**: desired output consists of one or more **continuous** variables
- **Classification**: Desired output consists of a finite number of **discrete** categories.

## (Linear) Classifier

- Goal:
  - Take an input vector  $\mathbf{x}$  and assign it to one of  $K$  discrete classes  $y$ .
- Assume an partition of the feature space:  $y=f(\mathbf{x})$
- Given examples  $(\mathbf{x}_i, y_i)$ , which may be noisy
- Learn  $f(\mathbf{x})$ , to enable prediction of  $y^*$  given new point  $\mathbf{x}^*$ . It should **generalise** well to new  $\mathbf{x}^*$ 
  - E.g.,  $x_1$ : Fish Weight,  $x_2$ =Fish Length,  $y$ =Fish Species.



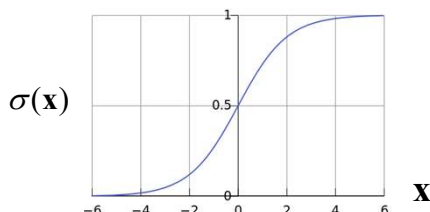
## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions

## Linear Versus Logistic Regression

- Linear & Logistic Regression use different representation/model assumptions.
- Linear Regression  $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \in [-\infty, \infty]$

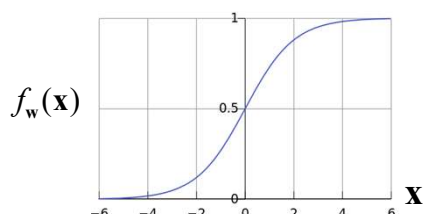
- Logistic Regression  $f_{\mathbf{w}}(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})} \in [0, 1]$



Sigmoid aka  
Logistic  
Function  $\sigma(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{x}})}$

## Linear Versus Logistic Regression

- Think about this Logistic function.
  - What if  $x \rightarrow -/+ \text{infinity}$ ? (Assume  $w=1$ )
    - $x$  controls range of  $\sigma(x)$  in  $[0,1]$
  - What if  $x=2$  and we increase/decrease  $w$  magnitude?
    - $w$  Controls slope of  $f_w(x)$



$$f_w(\mathbf{x}) = \frac{1}{(1 + e^{-w^T \mathbf{x}})}$$

Sigmoid aka  
Logistic  
Function  $\sigma(\mathbf{x}) = \frac{1}{(1 + e^{-x})}$

## Linear Versus Logistic Regression

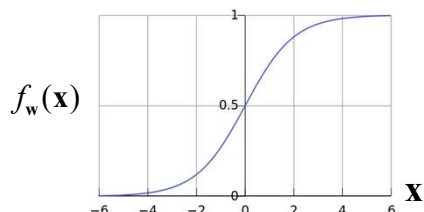
- Linear & Logistic Regression use different representation/model assumptions.
- **Linear** Regression  $f_w(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$

Range:  
+inf to  
- inf

- **Logistic** Regression

$$f_w(\mathbf{x}) = \frac{1}{(1 + e^{-w^T \mathbf{x}})}$$

Range: 0 to 1



Sigmoid aka  
Logistic  
Function  $\sigma(\mathbf{x}) = \frac{1}{(1 + e^{-x})}$

## Logistic Regression: What does it mean for modeling classes?

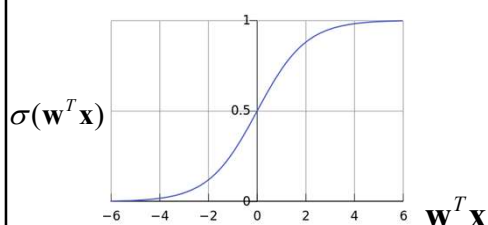
- With assumption

$$f_{\mathbf{w}}(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

- Probability Model:

$$p(y | \mathbf{x}, \mathbf{w}) = \text{Bernoulli}(y | \sigma(\mathbf{w}^T \mathbf{x}))$$

— => Probability of class  $y$  is  $\sigma(\mathbf{w}^T \mathbf{x})$



## Logistic Regression

$$p(y | \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

Range: 0 to 1

$$P(\text{Class 1} | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$$

$$P(\text{Class 0} | \mathbf{x}) = 1 - \sigma(\mathbf{w}^T \mathbf{x})$$

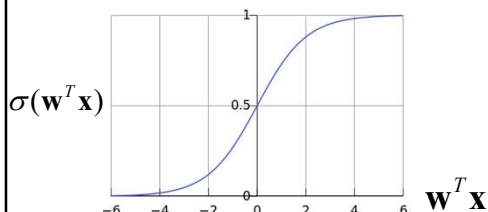
$$\sigma(\mathbf{w}^T \mathbf{x}) > 0.5 : \text{Class 1}$$

$$\sigma(\mathbf{w}^T \mathbf{x}) < 0.5 : \text{Class 0}$$

$$\mathbf{w}^T \mathbf{x} > 0$$

$$\mathbf{w}^T \mathbf{x} < 0$$

$$\mathbf{w}^T \mathbf{x} = 0$$



## What Kind of Decision Boundary Does LR Have?

- Boundary between two classes. Contour where:
  - $p(\text{Class 1}|\mathbf{x})=p(\text{Class 0}|\mathbf{x})=0.5$ .

$$1 / (1 + e^{-\mathbf{w}^T \mathbf{x}}) = 0.5$$

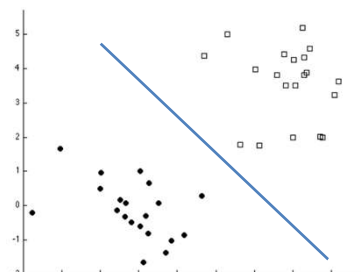
$$2 = 1 + e^{-\mathbf{w}^T \mathbf{x}}$$

$$0 = \mathbf{w}^T \mathbf{x}$$

=> It's a straight line!

$$f_{\mathbf{w}}(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

Parameters  $\mathbf{w}$  specify exactly which straight line we try to separate the data with

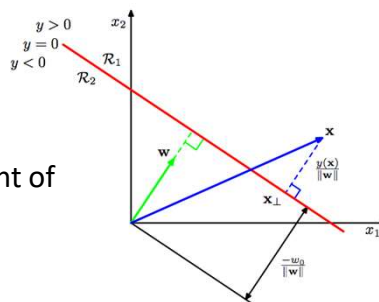


## Unpacking the decision boundary.

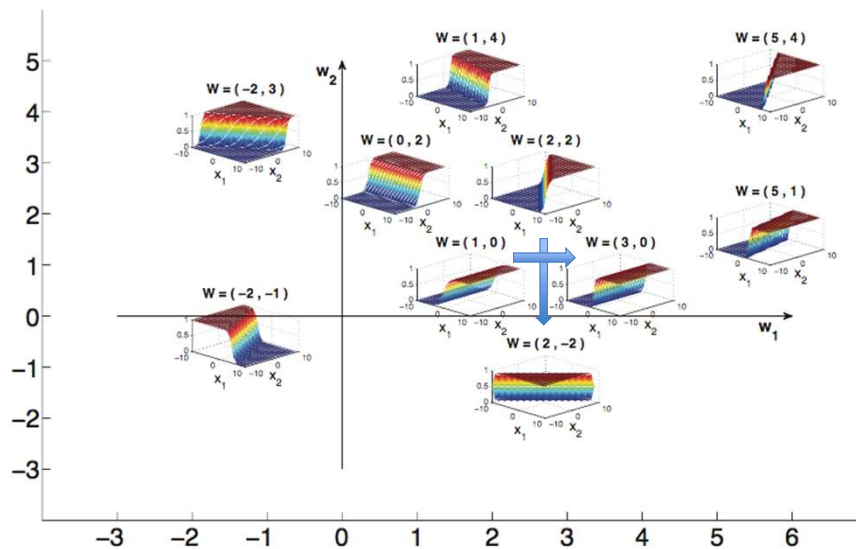
- Really we should be doing:  $\sigma(\mathbf{w}^T \mathbf{x} + w_0) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x} + w_0})}$

- Recall in linear regression, we used the trick  $\mathbf{x}' := [1, \mathbf{x}]$
- Now  $\mathbf{w}'$  automatically includes a  $[w_0, \mathbf{w}]$ .
- So  $\mathbf{w}'^T \mathbf{x}' = \mathbf{w}^T \mathbf{x} + w_0$

- The decision boundary (red) is perpendicular to the vector  $\mathbf{w}$ .
  - Weight  $w_0$  controls the displacement of the line from origin.



## Unpacking the prediction surface.



## Choosing the objective function

- For Linear Regression, we had:

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \quad \Rightarrow \quad E(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- For Logistic Regression?

$$f_{\mathbf{w}}(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})} \quad \Rightarrow \quad ??$$

## Choosing the objective function

- For Linear Regression, we had:

Square deviation  
of prediction from  
label

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \quad \Rightarrow \quad E(\mathbf{w}) = \sum_i (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

- For Logistic Regression?

Recall  $y_i \in \{0,1\}$

$$p(Y | X, \mathbf{w}) = \prod_i p(y = 1 | \mathbf{x}_i)^{y_i} (1 - p(y = 1 | \mathbf{x}_i))^{(1-y_i)}$$

$$f_{\mathbf{w}}(\mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

Probability model  $\mathbf{w}$   
would have got  
each point correct

Cost:  
Negative  
Log Likelihood

$$E(\mathbf{w}) = - \sum_i y_i \log p(y = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y = 1 | \mathbf{x}_i))$$

## Learning Logistic Regression

- Find the weights  $\mathbf{w}$  that minimize the cost function for data  $\{\mathbf{y}, \mathbf{X}\}$  or  $\{(y_i, \mathbf{x}_i)\}$ :

$$p(y = 1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

$$E(\mathbf{w}) = - \sum_i y_i \log p(y = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y = 1 | \mathbf{x}_i))$$

– Derivates wrt weights...

$$\frac{dE(\mathbf{w})}{d\mathbf{w}} = - \sum_i (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i$$

$$\frac{dE(\mathbf{w})}{d\mathbf{w}} = -X^T (p(\mathbf{y} | X) - \mathbf{y})$$



## Aside: Convex vs Closed Form

- Find the weights  $\mathbf{w}$  to minimize the cost function for data  $\{\mathbf{y}, \mathbf{X}\}$  or  $\{y_i, \mathbf{x}_i\}$ :

$$p(y=1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}}$$

$$E(\mathbf{w}) = -\sum_i y_i \log p(y=1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y=1 | \mathbf{x}_i))$$

– Derivates wrt weights...

$$\frac{dE(\mathbf{w})}{d\mathbf{w}} = -\sum_i (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i$$

Note: This time (unlike linear regression), we can't re-arrange to solve for  $\mathbf{w}$ ! It's stuck inside the sigmoid.



No closed form solution. Only gradient. However it is convex. Remember what convex means?

Gradient will get the solution as unique minimum.

## Learning Logistic Regression

- Find the weights  $\mathbf{w}$  that minimize the cost function for data  $\{\mathbf{y}, \mathbf{X}\}$  or  $\{y_i, \mathbf{x}_i\}$ :

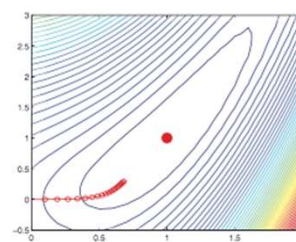
$$E(\mathbf{w}) = -\sum_i y_i \log p(y=1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y=1 | \mathbf{x}_i))$$

- Algorithm:

– Repeat:

$$\mathbf{w}' := \mathbf{w} + \alpha \frac{dE(\mathbf{w})}{d\mathbf{w}} \quad \mathbf{w}' := \mathbf{w} - \alpha \sum_i (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i$$

Until convergence



## Checkpoint

- Good practice:
  - Check you understand the **dimensions** of any algorithm/linear algebra expression.
- Questions:
  - If input data is  $d$  dimensions, how many dimensions is  $\mathbf{w}$ ?  $\mathbf{x}$ ?
  - How many dimensions is  $\sigma(\mathbf{w}^T \mathbf{x})$
  - How many dimensions is  $E(\mathbf{w})$ ?
  - How many dimensions is  $dE(\mathbf{w})/d\mathbf{w}$ ?

$$p(y=1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

$$E(\mathbf{w}) = -\sum_i y_i \log p(y=1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y=1 | \mathbf{x}_i))$$

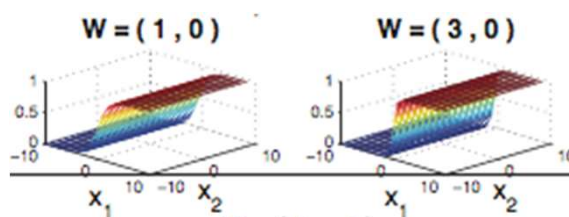
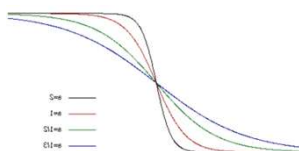
$$\frac{dE(\mathbf{w})}{d\mathbf{w}} = -\sum_i (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i$$

## What about Confidence?

- Larger magnitude weights result to steeper sigmoid.

$$p(Y | X, \mathbf{w}) = \prod_i p(y=1 | \mathbf{x}_i)^{y_i} (1 - p(y=1 | \mathbf{x}_i))^{(1-y_i)}$$

$$p(y=1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$



## What about Confidence?

- Larger magnitude weights result to steeper sigmoid.

$$p(Y | X, \mathbf{w}) = \prod_i p(y = 1 | \mathbf{x}_i)^{y_i} (1 - p(y = 1 | \mathbf{x}_i))^{(1-y_i)}$$

$$p(y = 1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

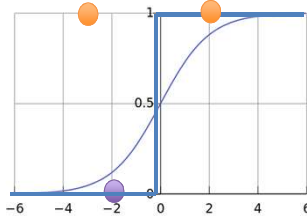
$$\text{— If } w=1. \sigma(\mathbf{w}^T \mathbf{x}_i) = [0.05, 0.9, 0.9]$$

$$\bullet \text{ Probability of all data} = 0.05 * 0.9 * 0.9 = 0.04$$

Learning won't increase weight

$$\text{— If } w=\text{inf. } \sigma(\mathbf{w}^T \mathbf{x}_i) = [0, 1, 1]$$

$$\text{Probability of all data} = 0 * 1 * 1 = 0.$$



Larger weights => more confident prediction.

## What about Confidence?

- Larger magnitude weights result to steeper sigmoid.

$$p(Y | X, \mathbf{w}) = \prod_i p(y = 1 | \mathbf{x}_i)^{y_i} (1 - p(y = 1 | \mathbf{x}_i))^{(1-y_i)}$$

$$p(y = 1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

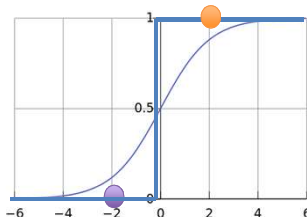
$$\text{— If } w=1. \sigma(\mathbf{w}^T \mathbf{x}_i) = [0.9, 0.9]$$

$$\bullet \text{ Probability of all data} = 0.9 * 0.9 = 0.8$$

$$\text{— If } w=\text{inf. } \sigma(\mathbf{w}^T \mathbf{x}_i) = [1, 1]$$

$$\text{Probability of all data} = 1 * 1 = 1.$$

Learning will increase weight forever



Larger weights => more confident prediction.

## Regularized Logistic Regression

- To address overconfidence, we can add an L2 norm regularizer.

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

## Aside: What is “l2 norm” about?

- Roughly: A **norm** gives a **scalar magnitude** from a **vector**.
  - What are ways you can imagine to do this?

$$\|\mathbf{w}\|_2 = \sqrt{\sum_k w_k^2} \quad \|\mathbf{w}\|_1 = \sum_k |w_k| \quad \|\mathbf{w}\|_0 = \sum_k w_k \neq 0$$

L2 norm  
“Euclidean”

L1 norm  
“Manhattan”

L0 norm

## Regularized Logistic Regression

- To address overconfidence, we can add a l2 regularizer.

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

- Cost is scalar.
  - Need a scalar to quantify how costly a weight vector is.
  - Need a norm of the weight vector.  $\lambda \sum_k w_k^2 = \lambda \mathbf{w}^T \mathbf{w}$

## Regularized Logistic Regression

- To address overconfidence, we can add a l2 regularizer.
  - As before we saw this corresponds to zero mean Gaussian weight prior!

$$\begin{aligned} p(\mathbf{w} | Y, X) &\propto p(Y | X, \mathbf{w}) p(\mathbf{w}) \\ &\propto \prod_i p(y_i = 1 | \mathbf{x}_i)^{y_i} (1 - p(y_i = 1 | \mathbf{x}_i))^{(1-y_i)} \exp(-\lambda \mathbf{w}^T \mathbf{w}) \end{aligned}$$

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

## Learning Regularized Logistic Regression

- Take derivatives of...
  - The Posterior rather than the Likelihood
  - =The augmented cost function

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

$$\mathbf{w}' := \mathbf{w} + \alpha \frac{E(\mathbf{w})}{d\mathbf{w}} \quad \mathbf{w}' := \mathbf{w} - \alpha \sum_i (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i - \lambda \mathbf{w}$$

## Summary

- Logistic Regression
  - Uses sigmoid function to predict probability in [0,1]
  - Weights specify the decision boundary.
  - Weight magnitudes specify its sharpness
- Cost as negative logprobability of correct assignment.
  - Using binary exponent trick.
- Differentiation wrt weights leads to convex gradient-based solution.
  - (But not closed form)
- Regularize to avoid overconfidence.

## Linear Classifier Examples – 1

### EECS Work 😊

- Undergraduate Project
  - $\mathbf{x}$ : Text of bills under debate in parliament
  - $y=f(\mathbf{x})=\mathbf{w}^T\mathbf{x}$ : Bill is accepted into law or not after parliament debate
- Simple linear classifier can predict with ~85% accuracy!



## Linear Classifier Examples – 2

### EECS Work 😊

- Person re-identification
  - Take two images (windows)  $x_1, x_2$
  - $F(x_1, x_2) = \mathbf{w}^T |x_1 - x_2|$  if it's the same person or not.



## Case Study – Feature Engineering

### EECS Work @ ACCV 2014 ☺

- Forensic Sketch / Caricature
  - Take an image  $\mathbf{x}$
  - $A_i(\mathbf{x}) = \mathbf{w}_i^T \mathbf{x}$  says if the face has an attribute  $i$  or not
    - E.g., big/small nose, fat/thin lips, long/short hair, male/female
    - Each  $A_i(\mathbf{x})$  is a new feature of original image
  - $F(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{w}^T |A(\mathbf{x}_1) - A(\mathbf{x}_2)|$  says if the same person or not.
  - => The feature itself is another machine learning model



Forensic sketch



Caricature

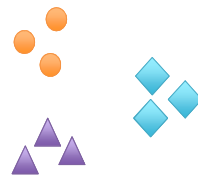
## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions



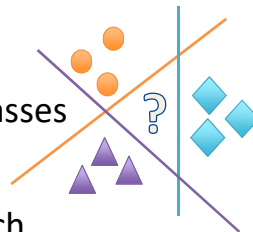
## From Binary to Multiclass

- We know how to model (logistic function) and learn (gradient descent) binary classifiers.
  - What about  $K > 2$ , multi-class problems?
  - Any ideas?



## Multiclass: 1 vs All

- Simple solution: 1-versus-All
  - Use  $K$  classifiers, each solving a two class problem of separating class  $k$  from all others.
- Issues:
  - Ambiguous regions
  - Gets expensive with many classes
- Next:
  - A 'natural' multiclass approach

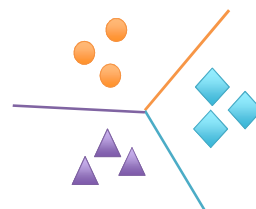


## Multiclass: Multinomial logistic regression, MaxEnt, softmax.

- Before we had  $p(y=1 | \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$   
 – ... a Bernoulli likelihood.  $= \text{Bernoulli}(y | \sigma(\mathbf{w}^T \mathbf{x}))$
- What was the “dice” rather than “coin” distribution?  
 – Multinomial. A parameter vector that adds up to 1.  
 – Now  $y=1\dots K$ , rather than  $y=0,1$ .

$$p(y | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_y^T \mathbf{x})}{\sum_{y'} \exp(\mathbf{w}_{y'}^T \mathbf{x})}$$

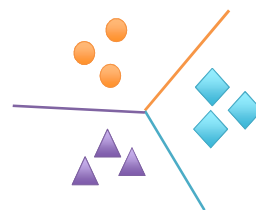
$$p(y | \mathbf{x}, \mathbf{W}) = \text{Multi}(y | \text{softmax}(\mathbf{w}_y^T \mathbf{x}))$$



## Think about Softmax...

- What's the range of  $\exp(\mathbf{w}^T \mathbf{x})$  for  $\mathbf{x} \Rightarrow -\text{inf}$  or  $+\text{inf}$ ??
- What's the range of:  $\frac{\exp(\mathbf{w}_y^T \mathbf{x})}{\sum_{y'} \exp(\mathbf{w}_{y'}^T \mathbf{x})}$
- What if  $\mathbf{x} = k\mathbf{w}_y$  for one  $y$ , and  $\mathbf{x} = -k\mathbf{w}_{y'}$  for the others  $y' \neq y$ ?  
 – And if  $k \Rightarrow \text{inf}$ ?

$$p(y | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_y^T \mathbf{x})}{\sum_{y'} \exp(\mathbf{w}_{y'}^T \mathbf{x})}$$



## Relating Softmax to Logistic

- For Binary case  $K = 2$ .

– What happens to Softmax?

$$p(y | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_y^T \mathbf{x})}{\sum_{y'} \exp(\mathbf{w}_{y'}^T \mathbf{x})}$$

$$p(y=1 | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_1^T \mathbf{x})}{\exp(\mathbf{w}_1^T \mathbf{x}) + \exp(\mathbf{w}_2^T \mathbf{x})} = \frac{1}{1 + \exp(\mathbf{w}_2^T \mathbf{x}) / \exp(\mathbf{w}_1^T \mathbf{x})}$$

$$= \frac{1}{1 + \exp(\mathbf{w}_2^T \mathbf{x} - \mathbf{w}_1^T \mathbf{x})} = \frac{1}{1 + \exp(-(\mathbf{w}_1 - \mathbf{w}_2)^T \mathbf{x})}$$

Same as  
Logistic From  
Before

$$p(y=1 | \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$

## Evaluating a Multiclass Classifier

- Quantify The Goodness of a MaxEnt Classifier?
  - Again, probability that all data assigned correctly.

$$p(Y | X, W) = \prod_i \prod_c p(c | \mathbf{x}_i, W)^{Y_{ic}}$$

– Where we let  $Y_{ic}$  be a sparse binary 1-of-N matrix encoding the labels. I.e., If  $y_i=c$ , then  $Y(i,c)=1$ .

– Dimensions of  $Y$  and  $W$ ?

- $Y$  is a  $N \times K$  matrix of labels.

0	0	1	0	0	0	...	0
0	0	0	1	0	0	...	0
0	0	1	0	0	0	...	0

- Now  $W$  is a  $K \times D$  matrix stacking  $D$  dims of weights for each of  $K$  classifiers  $\mathbf{w}_y$ .

## Evaluating a Multiclass Classifier

- Quantify The Goodness of a MaxEnt Classifier?
  - Again, probability that all data assigned correctly.

$$p(Y | X, W) = \prod_i \prod_c p(c | \mathbf{x}_i, \mathbf{w}_c)^{Y_{ic}}$$

- To specify a cost, take logs as usual.

$$E(\mathbf{W}) = \sum_i \sum_c Y_{ic} \log p(c | \mathbf{x}_i, W) \quad p(c | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_c^T \mathbf{x})}{\sum_{c'} \exp(\mathbf{w}_{c'}^T \mathbf{x})}$$

$$E(\mathbf{W}) = \sum_i \left[ \left( \sum_c Y_{ic} \mathbf{w}_c^T \mathbf{x}_i \right) - \log \left( \sum_{c'} \exp(\mathbf{w}_{c'}^T \mathbf{x}_i) \right) \right]$$

## Learning MaxEnt

- Differentiate the cost wrt W.

$$E(\mathbf{W}) = \sum_i \sum_c Y_{ic} \log p(c | \mathbf{x}_i, W)$$

$$p(c | \mathbf{x}, \mathbf{W}) = \frac{\exp(\mathbf{w}_c^T \mathbf{x})}{\sum_{c'} \exp(\mathbf{w}_{c'}^T \mathbf{x})}$$

$$E(\mathbf{W}) = \sum_i \left[ \left( \sum_c Y_{ic} \mathbf{w}_c^T \mathbf{x}_i \right) - \log \left( \sum_{c'} \exp(\mathbf{w}_{c'}^T \mathbf{x}_i) \right) \right]$$



$$\frac{dE(\mathbf{W})}{d\mathbf{w}_c} = \sum_i (p(c | \mathbf{x}_i, W) - Y_{ic}) \mathbf{x}_i$$

Intuition:

If  $\mathbf{x}_i$  supposed to be class  $c$ :  
 - Push  $\mathbf{w}_c$  towards/away  
 from  $\mathbf{x}$  as appropriate

(Differentiating the matrix one vector at a time)

## Summary

- Multi-class classification:
  - Simple indirect solution 1-vs-All of Logistic regressions
  - Direct solution: Use softmax.
- Cost function: Probability of correct assignment.
  - Use binary exponentiation trick again to express succinctly.
- Can learn with gradient as before.
- Same form as LR in the case where  $K=2$ .

## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions
  - Non-linear
  - Sparse regularizers
  - Multi-label

## Non-linear trick

- As with Linear Regression, we can use a fixed non-linear (e.g., polynomial) transformation.
  - Result: Non-linear in feature-space. **Linear wrt  $w$**

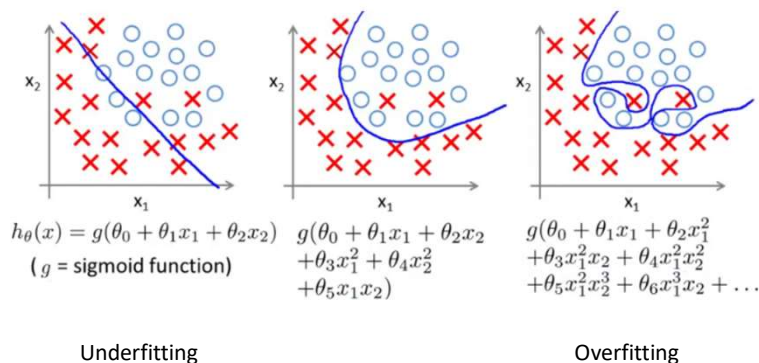
$$p(y=1|\mathbf{x}) = \sigma(\mathbf{w}^T \phi(\mathbf{x})) = \frac{1}{(1 + e^{-\mathbf{w}^T \phi(\mathbf{x})})}$$

$$\phi(\mathbf{x}) = [1, x_1, x_2, x_1^2, x_2^2, x_1 x_2]^T$$

$$p(y=1|\mathbf{x}) = \sigma(w_0 + w_1 x_1 + w_2 x_1^2 + w_3 x_2 + w_4 x_2^2 + w_5 x_1 x_2)$$

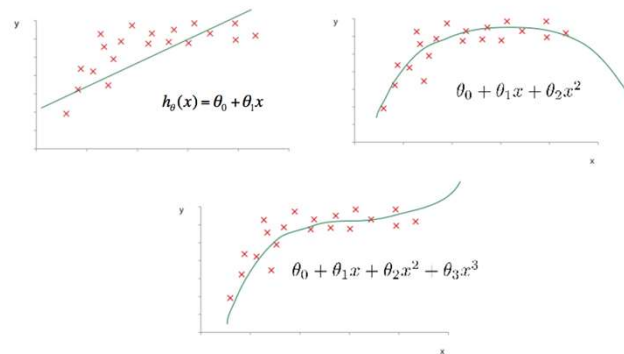
## Non-linear trick

- Using non-linear logistic regression this way, can result to over/underfitting again.



## Non-linear trick

- Contrast non-linear we saw in regression....



## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions
  - Non-linear
  - Sparse regularizers
  - Multi-label

## Regularization for Sparsity

- Common to use l2 regularization

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

- ...as for linear regression, l1 is also possible

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \sum_d |w_d|$$

- This increases **sparsity**...

## Regularization for Sparsity

- Common to use l2 regularization

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

- Sometimes we want sparsity.

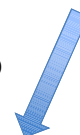
- Delete irrelevant distractor features.
- Save memory.
- Gain domain insight.

- How? Remember those regularizers...

- L0 is explicit about sparsity. Cost=# non-zero elements.

$$\|\mathbf{w}\|_2 = \sqrt{\sum_k w_k^2} \quad \|\mathbf{w}\|_1 = \sum_k |w_k| \quad \|\mathbf{w}\|_0 = \sum_k w_k \neq 0$$

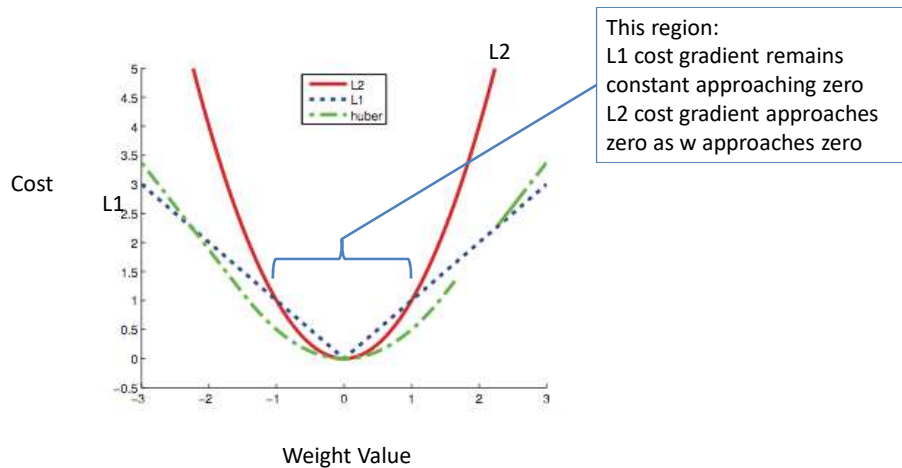
Can't  
differentiate  
wrt w





## L1 Regularization

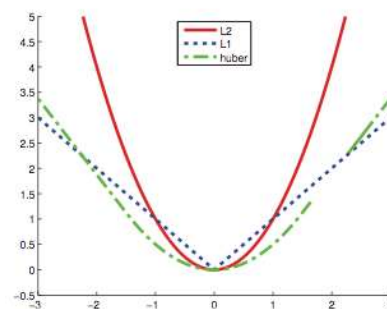
- L1 regularization increases sparsity.... Why?



## L1 Regularization

- But L1 regularisation still harder to Optimize. Why?
  - Absolute value has no derivative at zero (non-smooth corner)

$$\|\mathbf{w}\|_1 = \sum_k |w_k|$$



## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions
  - Non-linear
  - Sparse regularizers
  - Multi-label
  - Large Scale
  - Interpretation

## Multi-Label Classification

- Logistic Regression assumed data like  $\{\mathbf{x}_i, \mathbf{y}_i\}$ .
  - And trained predictor: 
$$p(y = 1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{(1 + e^{-\mathbf{w}^T \mathbf{x}})}$$
- What about if you have  $\{\mathbf{x}_i, \mathbf{y}_i\}$ ? (Vector of labels)
  - Classification analogy of multiple output linear regression.
  - When?
    - E.g., Webpage topic classifier for indexing. One web page can have multiple topics: Education+Government, Business+Finance, Business+Movie Industry..
    - Person can have multiple nationalities
    - Movie/music can span multiple genres.

## Multilabel vs Multiclass

- Multiclass
  - Predict  $p(\mathbf{y}|\mathbf{x})$ .
  - $y=1\dots K$ . Exclusive.
  - $\mathbf{x}$  has to have **exactly one** label.
- Multi-label
  - Predict  $p(\mathbf{y}|\mathbf{x})$
  - $\mathbf{y}$  is binary vector of size  $K$ .
  - If  $y_k=1$ , then  $\mathbf{x}$  has label  $k$ .
  - If  $y_k=1$  for all  $k$ . Then  $\mathbf{x}$  has all labels.
  - If  $y_k=0$  for all  $k$ . Then  $\mathbf{x}$  has no labels.

## Multilabel vs Multiclass

### Multiclass

- One  $\mathbf{x}$  has exactly one label
- Train data is scalar  $y$  for one  $\mathbf{x}$
- Prediction  $y=f(\mathbf{x})$  is a scalar
- $y=1\dots K$  exclusive
- $P(y|\mathbf{x})$  is a  $K$ -sized vector

### Multilabel

- One  $\mathbf{x}$  has many labels (0 to  $K$ )
- Train data is vector  $\mathbf{y}$  for one  $\mathbf{x}$ .
- Prediction  $\mathbf{y}=f(\mathbf{x})$  is a  $K$ -vector
- $y_k=\{0,1\}$
- $P(\mathbf{y}|\mathbf{x})$  is a  $K$ -sized vector.

## How to do Multi-label Classification?

- Given data  $\{\mathbf{y}_i, \mathbf{x}_i\}$ .
  - Train an independent logistic regression classifier for each  $p_k(y_k | \mathbf{x})$ .
  - Stack up the results  $p(\mathbf{y} | \mathbf{x}) = [p_1(y_1 | \mathbf{x}), \dots, p_K(y_K | \mathbf{x})]$ .
- This does not take into account label dependencies.

## Multi-label Example



## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions
  - Non-linear
  - Sparse regularizers
  - Multi-label
  - Large Scale
  - Interpretation

## Large Scale

- As for linear regression, the strategy is to minimise the cost function

$$E(\mathbf{w}) = -\sum_i y_i \log p(y_i = 1 | \mathbf{x}_i) + (1 - y_i) \log(1 - p(y_i = 1 | \mathbf{x}_i)) + \lambda \mathbf{w}^T \mathbf{w}$$

- By iterating:

$$\mathbf{w}^{s+1} := \mathbf{w}^s - \alpha \sum_i^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i - \lambda \mathbf{w}$$

- .... Means that you have to read database off disk into memory **very many times**.
  - Notice inner loop over  $i=1 \dots N$ 
    - ... within an iterative loop over  $s$

## Large Scale

- Regular:
  - Iterate S times:

$$\mathbf{w}^{s+1} := \mathbf{w}^s - \alpha \sum_i^N (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i - \lambda \mathbf{w}$$

- SGD with mini-batches:
  - Iterate M times:
  - Select a random subset B of total N examples
    - Iterate S times:

$$\mathbf{w}^{s+1} := \mathbf{w}^s - \alpha \sum_i^B (\sigma(\mathbf{w}^T \mathbf{x}_i) - y_i) \mathbf{x}_i - \lambda \mathbf{w}$$

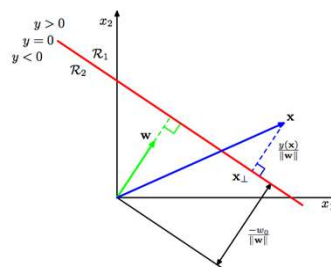
- Before: S memory reads of full database.
  - Now M\*(B/N) reads. Can be << S.
  - E.g., M=N/B => Read once.

## Overview

- Binary Classification – Logistic Regression
- Multiclass Classification – Max Entropy
- Extensions
  - Non-linear
  - Sparse regularizers
  - Multi-label
  - Large Scale
  - Interpretation

## Interpreting Logistic Regression Outputs

- Once learned:  $\mathbf{w} = \arg \min E(\mathbf{w})$ 
  - Recall each  $w_k$  is multiplied with  $x_k$  in  $\sigma(\mathbf{w}^T \mathbf{x})$
  - $\Rightarrow$  Positive  $w_k \Rightarrow$  Feature is indicator of class.
  - $\Rightarrow$  Negative  $w_k \Rightarrow$  Contra-indicator.
  - If l1 regularizer killed  $w_k$ 
    - Irrelevant to classification.



## LR Interpretation Example

EECS Work 😊

theguardian

all opinion culture business lifestyle fashion environment tech travel
≡ browse all sections

### Liberals swear more on Twitter than rightwingers, says study

Swearwords account for two of the top ten words used by left-leaning tweeters, while those on the right show a preference for "God" and "psalm"

An analysis of nearly a million tweets has found that liberals distinguish themselves from others online by letting rip with profanities.

While right-leaning tweeters displayed more caution in the language they used, the reds were more blue, as liberals peppered their 140-character missives with a hefty helping of "shits" and "fucks".

## LR Interpretation Example

EECS Work ☺

Table 7. Logistic regression model with all predictors using data without outliers.

Predictors	Estimate	Standard Error	Z value	P value
(Intercept)	-0.32546	0.22806	-1.427	0.153555
1 <sup>st</sup> person singular pronouns	0.08867	0.02054	4.316	1.59E-05***
1 <sup>st</sup> person plural pronouns	-0.3435	0.09793	-3.507	0.000452***
Swear words	0.88577	0.21676	4.086	4.38E-05***
Positive emotion words	0.13363	0.02517	5.31	1.10E-07***
Negative emotion words	-0.12299	0.05362	-2.294	0.021806*
Anxiety words	0.85417	0.21939	3.893	9.88E-05***
Feeling words	-0.03407	0.14776	-0.231	0.817626
Tentative words	-0.11298	0.05405	-2.09	0.036593*
Certainty words	-0.03662	0.07287	-0.503	0.615252
Achievement words	-0.09091	0.05773	-1.575	0.115288
Religion words	-0.22651	0.1444	-1.569	0.116749
Death words	-0.28486	0.24512	-1.162	0.245182

Similarly, study from colleagues at MSR....

### Private traits and attributes from digital records of human behavior

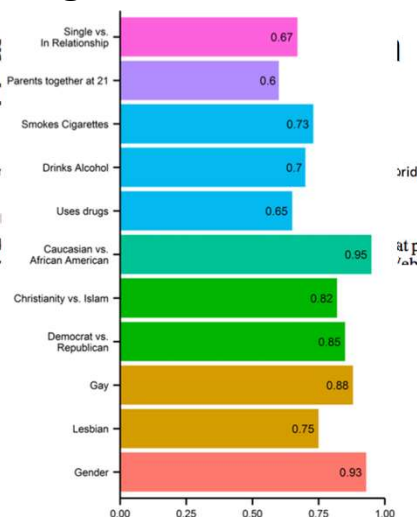
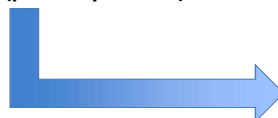
Michal Kosinski<sup>a,1</sup>, David Stillwell<sup>a</sup>, and Thore Graepel<sup>b</sup>

<sup>a</sup>Free School Lane, The Psychometrics Centre, University of Cambridge, Cambridge, United Kingdom

Edited by Kenneth Wachter, University of California, Berkeley, CA, and approved by the Royal Society

We show that easily accessible digital records of behavior, Facebook Likes, can be used to automatically and accurately predict a range of private traits and attributes.

Facebook Likes (public by default)





## Summary

- Non-linear extensions exist
  - Like linear regression, take fixed polynomial basis functions.
    - Then you would need to regularize for more than ‘just’ overconfidence.
- L2/L1 regularizer options also exist.
  - L2 is efficient.
  - L1 less efficient, but helps find sparsity.
- Multi-label versus Multi-class classification.
  - Multi-label: Simple solution by K-binary classifiers.